

The Secrets of UART FIFO

Casper Yang, Senior Product Manager

support@moxa.com

A UART (universal asynchronous receiver transmitter) is a key component of RS-232/422/485 serial communication hardware, and documents that introduce UARTs are readily available. A UART's FIFO buffer is designed to improve communication performance, although if it is configured incorrectly, you may find that your communication performance is degraded. In this paper we give a simple introduction to the UART FIFO, discuss how it influences communication behavior, and provide instructions on the proper way to configure a FIFO.

Why FIFO?

A FIFO (First In First Out) is a UART buffer that forces each byte of your serial communication to be passed on in the order received. For an 8250 or 16450 UART, for example, the FIFO has a size of only one byte. This means that the UART will issue an interrupt to the system for each byte of data received, which uses a lot of CPU resources. If you don't read the data from the UART in time, the next byte will overwrite the data. For higher baudrate applications (higher than 115200 bps, for example), a 1-byte limitation will make it more likely that data will be lost.

Copyright © 2009 Moxa Inc.

Released on Oct.01, 2009

About Moxa

Moxa manufactures one of the world's leading brands of device networking solutions. Products include serial boards, USB-to-serial hubs, media converters, device servers, embedded computers, Ethernet I/O servers, terminal servers, Modbus gateways, industrial switches, and Ethernet-to-fiber converters. Our products are key components of many networking applications, including industrial automation, manufacturing, POS, and medical treatment facilities.

How to Contact Moxa

Tel: 1-714-528-6777
Fax: 1-714-528-6778

Web: www.moxa.com
Email: info@moxa.com



This document was produced by the Moxa Technical Writing Center (TWC). Please send your comments or suggestions about this or other Moxa documents to twc@moxa.com.

The Trouble with FIFO—Timeout Design

The data loss that can result from a 1-byte limitation on FIFO size is why most advanced UARTs, such as the 16550A, support a FIFO size of 16 bytes or more. With a larger FIFO, for each interrupt that is issued you can often read all of the data at one time, which saves system CPU resources since you won't be reading data all the time. This is good for large data transfers, but may not be good for real time control. The frequency with which interrupts are issued can be controlled by adjusting the RTL (receive trigger level) value. When the amount of received data reaches the RTL, the UART will issue an interrupt. The 16550A with its 16-byte FIFO supports 4 RTL levels: 1, 4, 8, and 14. This means an interrupt will be issued when the UART FIFO receives 1, 4, 8, or 14 bytes of data. Note that the highest level is 14 instead of 16 to prevent overflow.

Another problem you need to consider is this: If the amount of received data hasn't reached the RTL value, when will the UART issue an interrupt? For example, if you configure the trigger level to 14 (the default), but the device only sends 4 bytes (substantially lower than 14), how long will you need to wait until the UART issues the receive interrupt?

A UART uses a timeout to solve this problem. The length of the timeout is different for different UART designs, but as a general rule, most UARTs wait the time needed to transmit 4 bytes. For example, for a serial configuration of 300 bps, 8 data bits, odd parity, and 2 stop bits, 1 bit needs 3.3 ms to be transferred. Since transmitting each byte requires transmitting a start bit, 8 data bits, 1 parity bit, and 2 stop bits, a total of 12 bits will be transmitted, and consequently:

$$4 \text{ bytes time} = 4 \text{ bytes} \times 12 \text{ bits/byte} \times 3.3 \text{ ms/bit} \approx 160 \text{ ms}$$

Such a long latency is not appropriate for some time critical control applications.

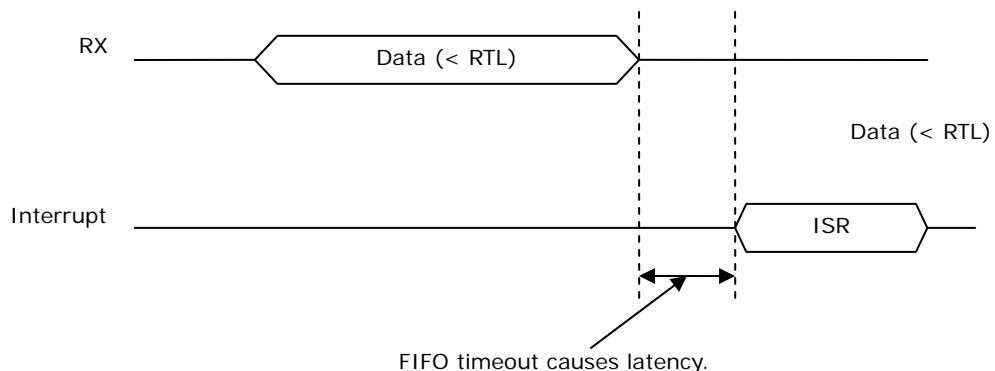


Fig. 2: FIFO Timeout Causes Latency

Receiving—Throughput or Latency?

In serial communication, there is a tradeoff between throughput and latency, and you can only choose one, which is fortunate since most applications only need one of the two. Applications such as downloading firmware or file downloads in CNC control only need good throughput. For such applications, you can simply configure the UART to the highest RTL value (for example, 14 bytes for the 16550A). But for other applications, you need to disable the FIFO or set the RTL value to 1 to get good (i.e., short) latency. This means that you need to respond as soon as possible when you receive data from the serial line. In software flow control for example, if you receive an XOFF character, you need to stop transmitting data immediately.

Transmitting—Slow Throughput with Lag for Slow Devices

The above discussion focuses on receiving data, but you also need to control how data is transmitted. The difference is in how many bytes are transmitted for each interrupt. Keep in mind that transmitting more bytes will result in a higher throughput. For some older devices with a small buffer, it is easy to get overrun in high throughput situations. The solution is to transmit one byte for each interrupt, which increases the lag between two bytes and gives the device more processing time.

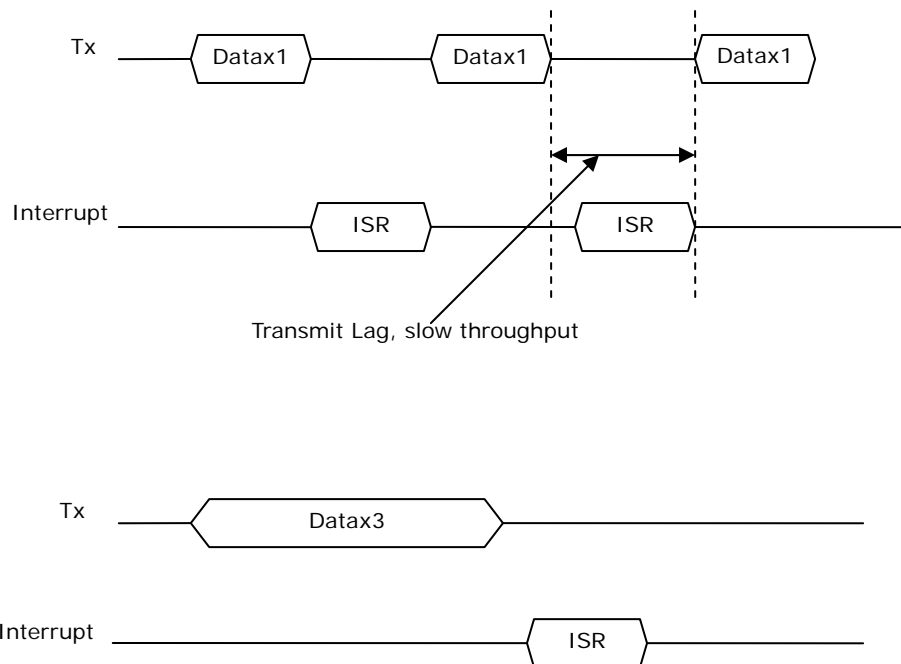


Fig. 3: A FIFO provides better Tx throughput, and fewer interrupts means using fewer CPU resources.

Better Receive Throughput—High/Low Water for Flow Control

Many advanced UARTs provide an on-chip flow control function that uses both a receive high water level (RBH) and low water level (RBL), as opposed to just one trigger level. If a UART only supports one receive trigger, it will issue an interrupt and turn off the RTS when the length of the data reaches a certain level, and it will turn on the RTS when all of the data has been read. At that moment there will be no data in the FIFO and the system will waste time waiting for additional data. For better throughput, set the RBL value to a non-zero number. For example, if you set RBH to 120 bytes and set RBL to 16 bytes for a UART with a 128-byte FIFO, the RTS will switch off when the received data length is greater than 120 bytes. The RTS will turn on again when fewer than 16 bytes are in the FIFO after the system moves some data to memory. The UART also provides an interrupt level (RBI) setting, which is equal to a traditional trigger level. For this UART, this means that the timing for issuing an interrupt and turning off the RTS can be different (note that traditional UARTs only use the RTL value).

Better Transmit Throughput—High/Low Water

Transmitting requires a similar design. Traditionally, the UART would issue a Tx Empty interrupt (TxINT) when all of the data in the FIFO was sent out. At that moment, the UART wasted time waiting for data. Many advanced UARTs provide a low water level (TBL) setting to avoid this problem. You can set the TBL value to a nonzero number to eliminate the transmit lag and get higher throughput. But you need to use it carefully. For some RS-485 half duplex applications, you need to control RTS for switching between Tx and Rx. This means that you need to turn on the RTS before transmitting data and turn off the RTS after all of the data has been sent out. Most drivers will turn off the RTS when a Tx Empty interrupt is issued. If any data is still located in the FIFO, turning off the RTS will cause a Tx problem. To get around this, just set the low water level to zero to make sure there is no data in the FIFO when the Tx Empty interrupt is issued.




Fig. 4: With a low water setting, you can put data in the ISR before the FIFO is empty to eliminate Tx lag.

Conclusion

In this paper we introduce a number of issues related to UART FIFOs. A FIFO is good for improving throughput, but it can also cause problems, even if you use the default settings. Knowing your application in detail can prevent problems and improve throughput, and result in better communication with your serial devices.

FAQ	Question	Answer
1	How can I get the best Rx latency?	Disable the FIFO, or set the Rx trigger level to 1.
2	If my serial device is slow and has a small buffer, what can I do to prevent data overflow?	Disable the FIFO, or set RTL=1 and send one byte for each interrupt.
3	Is there any way to get better Rx throughput or eliminate Rx lag?	Use an advanced UART with RBH and RBL, and set the RBL to a nonzero value.
4	Is there any way to get better Tx throughput or eliminate the Tx lag?	Use an advanced UART with TBH and TBL, and set TBL to a nonzero value.
5	What is the proper FIFO setting when working with RS-485 with RTS flow control?	Set TBL to zero if using an advanced UART.
6	I'm using XON/XOFF flow control but am getting data overflow. How can I solve this problem?	For a UART that does not have on-chip XON/XOFF flow control, try disabling the FIFO.